

The Evidence Trap?

Showing that "fact-based" Conflicts Never Settle

Ryan Healy
Burlington, Vermont
rah5ff@virginia.edu

May 2026

"I do not paint things, I paint only the differences between things."

Henri Matisse, 1943

THE DATA AND THE QUESTION

THIS PROJECT draws on my [Capstone research](#) into Wikipedia's Arbitration Committee (or ArbCom). This community elected body serves as English Wikipedia's court of last resort. A case only reaches ArbCom after every lower-level mechanism (talk-page negotiation, third opinion, mediation, admin intervention) has already failed. That failure is the *boundary*, a case going to ArbCom is something the Wikipedia community struggles to resolve.

The data set covers 53 cases spanning 2004–2025, organized into six conflict categories: **Politics, Religion, Science, Social, Geopolitics, and Policy**. Each case carries timestamped edit events, phase-level breakdowns (evidence, workshop, proposed decision, outcome), and final dispositions. Most ArbCom cases tend to deal with discipline cases. This sample was selected to instead be about hard to resolve topics instead of individual bans and appeals.

The question is *Why do some arbitration cases never really close?*

A conduct dispute ends when the disruptive editor is removed. A factual dispute should end with the facts of the case, but some questions either get stuck "waiting for more

data" or are purposefully left open. Some categories of disputes are different, how does *Religion* and *Science* get stuck always looking for data but *Politics* is better at resolving the facts of the case?

THE VISUALIZATION

The file `evidence_trap.html` encodes each case as a binned 60-segment horizontal row. The key encoding decision was to strip everything back to a single channel. Here, the display is reorganized to highlight only one color: **orange evidence activity**. The brightness of each orange cell scales with the fraction of edits in that window that belong to the evidence phase.

For this viz, I Striped the visualization back to a single encoded channel with evidence fraction, expressed as orange brightness. This chart makes one argument and one argument only, and the group-average rows allow the reader to see the whole claim in five seconds without opening a single individual case. A faint glow means that the case brushed the evidence and moved on; a deep amber means that the case is still building the record at that stage of its timeline.

The cases are grouped into six labeled bands. Each band leads with a **group-average row**: a single bar showing the mean evidence fraction across all cases in the category, smoothed across the 60 windows. Individual case rows expand below it on click. The group average is the primary reading surface; individual cases provide ground truth for what drives the aggregate. I included some notes about the cases for the curious reader that clicks into the graphic for more detail. The interactive version is published at observablehq.com/d/7bff75758de5a3c9; the full visualization is embedded on the following page.

05/05/2026, 10:58

The Evidence Trap · Wikipedia Arbitration



Figure 1: Interactive visualization of all 53 ArbCom cases grouped by dispute type (live version). Each row is a case's **group average** evidence fraction across 60 equal time windows. **Orange** intensity = share of edits belonging to the evidence phase; dark cells = no evidence activity. **Fact-based groups** (Geopolitics, Science, Religion): orange persists well past the midpoint — the evidence record is never considered complete. **Conduct-based groups** (Politics, Policy, Social): orange concentrates in the first third then clears — the committee gathers what it needs, rules, and closes. Rows expand on click to reveal individual cases.

FINDINGS

The six group-average rows tell the story without requiring any individual case to be opened.

Politics and Policy clear evidence fast. Policy cases follow the same grammar: the evidence phase is bounded, the committee deliberates, and a procedural ruling ends the proceedings cleanly. Both groups deal in conduct and rule-application, questions with answers the committee can actually deliver.

Science and Geopolitics never release the evidence phase. The orange glow in the Science and Geopolitics group-average bars does not dim after the first third of the timeline. It persists at mid-to-high brightness well past the midpoint and in some cases runs all the way to the end. Climate Change returned to arbitration three times in five years. Palestine–Israel reached the committee five times over two decades. Each time, the evidence phase re-ignited because the underlying factual dispute (what happened, which source is reliable, whose framing is neutral) was not resolved by the previous ruling.

The bottleneck is not the ruling. It is building the record. The proposed-decision band is almost always thin across all six groups. Committees decide quickly once they act. What the orange bars reveal is that in fact-based disputes, “acting” is perpetually deferred because the evidence record is never considered complete. New editors arrive, challenge previous sourcing conclusions, and the orange phase reactivates.

REFLECTION

I went through several iterations to come up with this picture. Originally I wanted to focus on the edits and revert “wars” for some of the longer term topics. The difficulty ended up in the “so what”, and it was too easy to have the reader get caught in the specifics of each case rather than seeing the structure. You don’t need to know the difference between “work-shopping” and “evidence” or that some cases end up snowballing and being about way more than they started on. I had trouble teasing out that pattern, so I did what I could to really distill down to the “evidence” of each case.

Two things I would improve.

1) the group-average row uses a simple arithmetic mean across all 60 windows, which can obscure the bimodal signal visible in Religion: some cases clear the evidence phase early and some never do, and those two patterns average into a flat middle that undersells the actual spread. A small density strip alongside the mean bar would show whether the group average is tight or hiding two sub-populations.

2) There doesn’t seem to be a real sense of scale of editor involvement or how long some of these things took. This was by design, as they were distracting from the main point, but there is something there that should be revisited.

Thank you for a great Semester! I look forward to your feedback.

AI DISCLOSURE

I did help use the Claude code to help draft the code for this viz. All prompts and data are available on request.

Repositories:

- **Capstone pipeline** (data collection & processing): github.com/rahds/Wikipedia_Dispute_Models (private – email rah5ff@virginia.edu to request access)
- **Data visualization coursework** (this assignment & all visualizations): github.com/rahds/Data_Viz_Class
- **Interactive visualization** (Observable notebook): observablehq.com/d/7bff75758de5a3c9